

```
In [8]: # Define a threshold value for anomalies
threshold_lstm = np.percentile(mse_lstm, 95)

# Predict anomalies
y_pred_lstm = mse_lstm > threshold_lstm

# Ensure y_pred_lstm has the same length as y_test
y_pred_lstm = y_pred_lstm[:len(y_test)]

# Flatten predictions for evaluation
y_pred_lstm = y_pred_lstm.flatten()

print(model_lstm.summary())

# Evaluate the performance
from sklearn.metrics import classification_report
print("LSTM Model Performance")
print(classification_report(y_test, y_pred_lstm))
```

Model: "LSTM_AnomalyDetectionModel"

Layer (type)	Output Shape	Param #
LSTM_Layer_1 (LSTM)	(None, 1, 64)	24,320
Dropout_Layer_1 (Dropout)	(None, 1, 64)	0
LSTM_Layer_2 (LSTM)	(None, 32)	12,416
Dropout_Layer_2 (Dropout)	(None, 32)	0
Output_Layer (Dense)	(None, 30)	990

Total params: 113,180 (442.11 KB)

Trainable params: 37,726 (147.37 KB)

Non-trainable params: 0 (0.00 B)

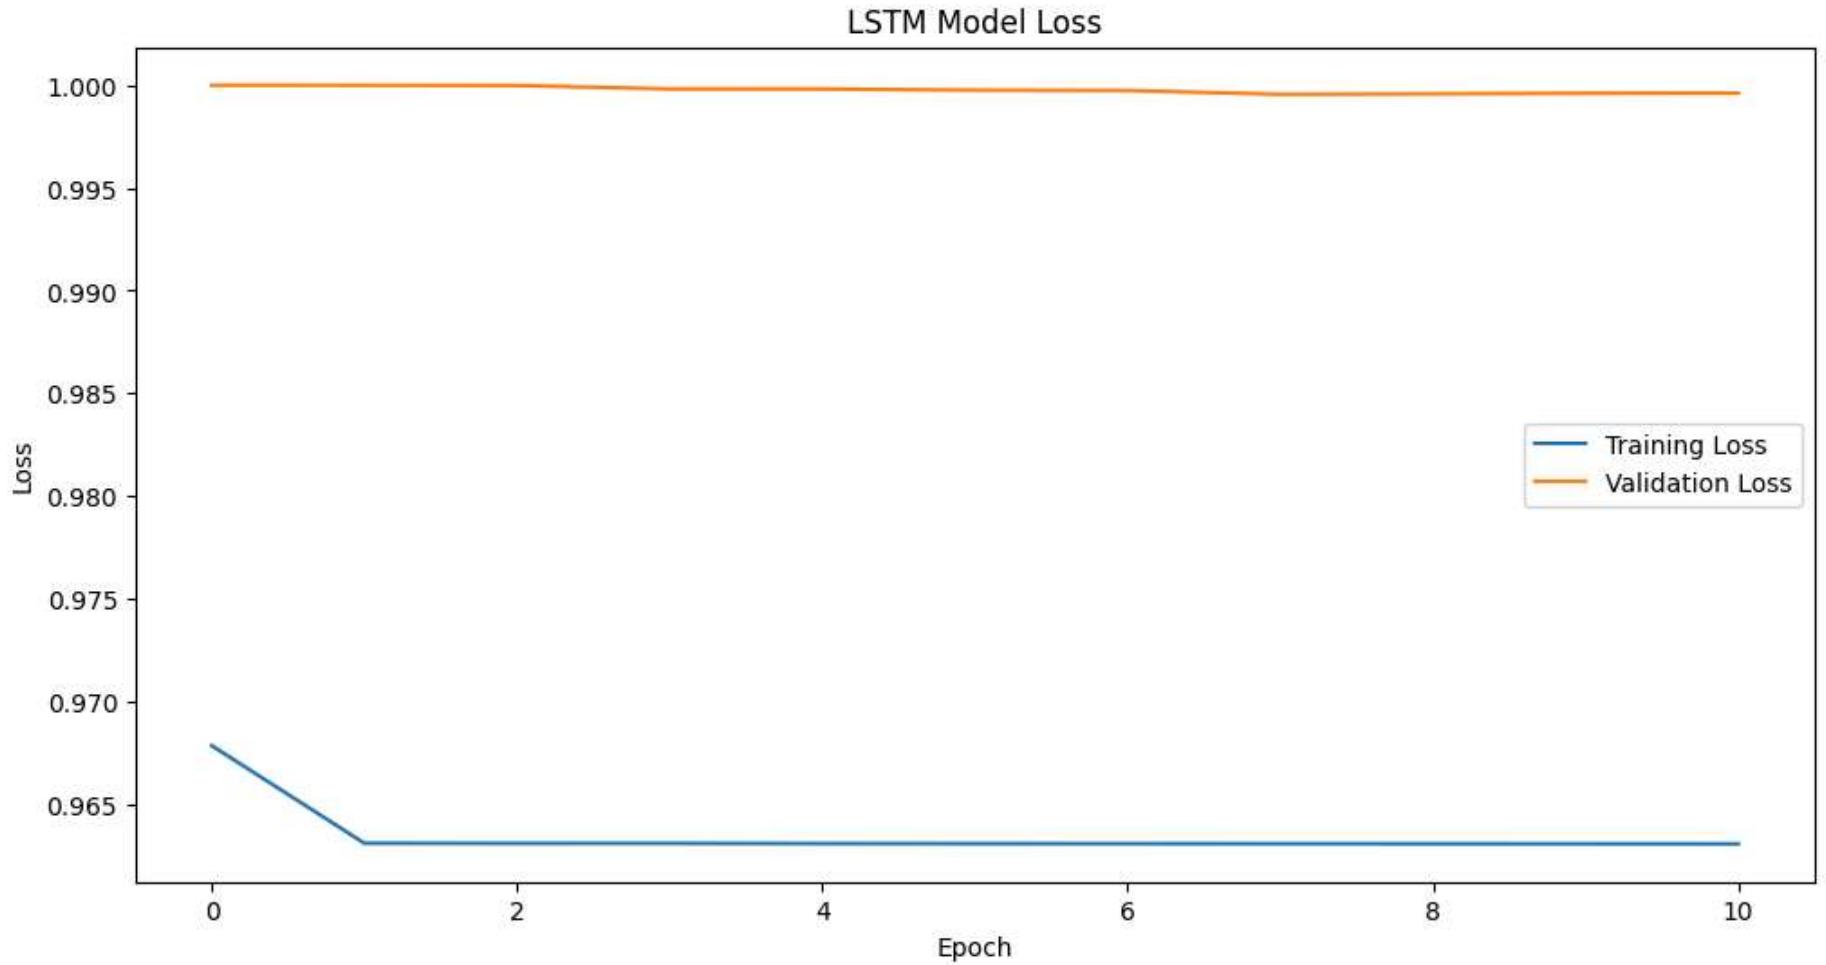
Optimizer params: 75,454 (294.75 KB)

```
None
LSTM Model Performance
precision    recall   f1-score   support

          0      1.00      0.94      0.97    284315
          1      0.00      0.08      0.00      492

   accuracy         0.94    284807
macro avg       0.50      0.51      0.49    284807
weighted avg     1.00      0.94      0.97    284807
```

```
In [9]: # Plot LSTM model Loss
plt.figure(figsize=(12, 6))
plt.plot(history_lstm.history['loss'], label='Training Loss')
plt.plot(history_lstm.history['val_loss'], label='Validation Loss')
plt.title('LSTM Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



Recurrent Neural Networks (RNN)

- **Core Functionality:** Basic sequential pattern recognition
- **Comparative Characteristics:**
 - Simpler architectural design
 - Faster computational processing
 - Limited long-term dependency capture

Building the RNN Model Architecture

Layer Breakdown

- **Input Layer** (SimpleRNN with 64 units):
 - **Purpose:** This layer is responsible for reading the input features and capturing the dependencies over time.
 - **64 Units:** The number of units determines the network's capacity to learn complex patterns. We chose 64 units to balance model complexity and performance.
 - **return_sequences=True:** This ensures that the output from each time step is fed to the next layer, preserving the temporal sequences.
- **Dropout Layer** (0.2):
 - **Purpose:** Dropout helps prevent overfitting by randomly setting a fraction of the input units to 0 at each update during training time.
 - **0.2:** This means 20% of the units will be dropped randomly, promoting generalization in the model.
- **Second SimpleRNN Layer** (32 units):
 - **Purpose:** This layer further processes the output from the previous RNN layer to capture higher-level temporal patterns.
 - **32 Units:** A smaller number of units are used as we move deeper into the network, focusing on refining and condensing information.
 - **return_sequences=False:** Since this is the final RNN layer, we set it to return only the last output in the sequence, which is then passed to the dense output layer.
- **Second Dropout Layer** (0.2):
 - **Purpose:** This dropout layer continues to prevent overfitting by dropping 20% of the units, further ensuring that the model does not rely too heavily on specific neurons.
- **Output Layer** (Dense with Sigmoid Activation):
 - **Purpose:** The final dense layer with a single unit outputs the probability of a transaction being fraudulent.
 - **Sigmoid Activation:** The sigmoid function maps the output to a value between 0 and 1, making it suitable for binary classification.

Summary

- **Stacked RNN Layers:** The combination of stacked SimpleRNN layers helps the model learn both immediate and long-term dependencies in transaction sequences.
- **Dropout Layers:** These are crucial for preventing overfitting, especially when working with sequential data, ensuring that the model generalizes well to unseen data.
- **Output Layer:** The dense layer with sigmoid activation provides the final binary classification decision.

```
In [18]: #STEP 1: Import Libraries
# Already done

#STEP 2: Preprocess the Data
# nothing to do, we reuse the same data as for LSTM model also for the RNN model

#STEP3: Build the RNN Model
# Define the RNN model
model_rnn = Sequential(name='RNN_AnomalyDetectionModel')
model_rnn.add(Input(shape=(1, X_train_reshaped.shape[2]), name='Input_Layer'))
model_rnn.add(SimpleRNN(64, return_sequences=True, name='RNN_Layer_1'))
model_rnn.add(Dropout(0.2, name='Dropout_Layer_1'))
model_rnn.add(SimpleRNN(32, return_sequences=False, name='RNN_Layer_2'))
model_rnn.add(Dropout(0.2, name='Dropout_Layer_2'))
model_rnn.add(Dense(X_train_reshaped.shape[2], activation='sigmoid', name='Output_Layer'))

#STEP4: Train the Model
# Train the LSTM model
model_rnn.compile(loss='mean_squared_error', optimizer=Adam(learning_rate=0.001))
early_stopping = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
history_rnn = model_rnn.fit(X_train_reshaped, X_train_reshaped, epochs=20, batch_size=64, validation_data=(X_test_reshaped, X_test_reshaped))
print(model_rnn.summary())

#STEP5: Evaluate the Model
# Calculate reconstruction error for RNN
# Evaluate in Batches
batch_size = 1000
mse_rnn = np.array([], dtype=np.float64)
mse_rnn = np.array([])
for start in range(0, X_test_reshaped.shape[0], batch_size):
    end = min(start + batch_size, X_test_reshaped.shape[0])
    reconstructed_rnn_batch = model_rnn.predict(X_test_reshaped[start:end])
    mse_rnn_batch = np.mean(np.power(X_test_reshaped[start:end]-reconstructed_rnn_batch, 2), axis=2).flatten() # Ensure mse_rnn_batch is a 1D array
    mse_rnn = np.concatenate((mse_rnn, mse_rnn_batch), axis=0)
```

Epoch 1/20
4443/4443 27s 5ms/step - loss: 0.9805 - val_loss: 1.0000

Epoch 2/20
4443/4443 21s 5ms/step - loss: 0.9419 - val_loss: 0.9999

Epoch 3/20
4443/4443 20s 5ms/step - loss: 0.9675 - val_loss: 0.9998

Epoch 4/20
4443/4443 21s 5ms/step - loss: 0.9570 - val_loss: 0.9998

Epoch 5/20
4443/4443 23s 5ms/step - loss: 0.9611 - val_loss: 0.9997

Epoch 6/20
4443/4443 25s 6ms/step - loss: 0.9742 - val_loss: 0.9997

Epoch 7/20
4443/4443 22s 5ms/step - loss: 0.9663 - val_loss: 0.9997

Epoch 8/20
4443/4443 24s 5ms/step - loss: 0.9556 - val_loss: 0.9997

Epoch 9/20
4443/4443 24s 5ms/step - loss: 0.9692 - val_loss: 0.9996

Epoch 10/20
4443/4443 26s 6ms/step - loss: 0.9652 - val_loss: 0.9996

Epoch 11/20
4443/4443 22s 5ms/step - loss: 0.9674 - val_loss: 0.9996

Epoch 12/20
4443/4443 23s 5ms/step - loss: 0.9661 - val_loss: 0.9996

Epoch 13/20
4443/4443 21s 5ms/step - loss: 0.9727 - val_loss: 0.9996

Epoch 14/20
4443/4443 25s 6ms/step - loss: 0.9620 - val_loss: 0.9996

Model: "RNN_AnomalyDetectionModel"

Layer (type)	Output Shape	Param #
RNN_Layer_1 (SimpleRNN)	(None, 1, 64)	6,080
Dropout_Layer_1 (Dropout)	(None, 1, 64)	0
RNN_Layer_2 (SimpleRNN)	(None, 32)	3,104
Dropout_Layer_2 (Dropout)	(None, 32)	0
Output_Layer (Dense)	(None, 30)	990

Total params: 30,524 (119.24 KB)

Trainable params: 10,174 (39.74 KB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 20,350 (79.50 KB)